

Distributed shared memory through key-value stores

Mantas Serapinas

4th Year Project Report
Artificial Intelligence and Computer Science
School of Informatics
University of Edinburgh

2018

Abstract

* Still to be written.

Acknowledgements

* Still to be written.

Table of Contents

1	Introduction	7
1.1	Motivation	7
1.2	Scope	8
1.3	Contributions	8
1.4	Synopsis	8
2	Data store for DSM	11
2.1	Overview	11
2.2	Bigtable	11
2.3	Datastore	12
2.4	Spanner	12
2.5	Benchmarking results	12
2.5.1	Loading the data	13
2.5.2	Workloads	13
2.5.3	Conclusions	15
2.6	Reading and writing the contents of Bigtable	15
2.7	Issues encountered	18
3	Calling get/put functions on Bigtable	19
3.1	Research on possible solutions	19
3.1.1	Source code level solutions	19
3.1.2	Compiler level solution	20
3.2	LLVM	20
3.3	Architecture	21
3.4	LLVM pass	23
4	Custom heap allocator	25
4.1	Motivation	25
4.2	Implementation	25
4.3	Consistency on multithreaded programs	28
5	Translation improvements	29
5.1	Different type of arguments to get and put functions	29
5.2	Detecting store and load instructions to heap memory	29
6	Results	31

6.1	Full and heap translation comparison	31
6.2	Comparison to uninstrumented program	31
7	Conclusions	33
7.1	Overview	33
7.2	Future work	33
7.2.1	Better support for structured data	33
	Bibliography	35

Chapter 1

Introduction

1.1 Motivation

Distributed shared memory (DSM) is memory architecture where physically distributed memory can be accessed as one logically shared address space. Systems based on shared memory architecture reduce the complexity of parallel programming (Z.Huang et al., 2006). Unfortunately, building an efficient distributed shared memory system is a huge challenge and the documentation on the existing open-source DSMs is rather limited. Thus it can be a daunting task to run parallel programs on distributed shared memory systems.

With cloud computing becoming increasingly popular new solution became available, namely NoSQL data stores (Grolinger et al., 2013). NoSQL can be completely schema-free, most popular data models being key-value stores, document stores, column-family stores, and graph databases. It is able to scale horizontally over many commodity servers. On top of that, some cloud data management systems provide strong consistency model, which means that after update operations all nodes agree on the new value before making it available to the user. All these properties make it possible to use such data stores as distributed shared memory.

The focus of this project is to expose the distributed shared memory model in a cloud by implementing an instrumentation tool which translates load and store instructions to get and put calls to key-value store. This tool will let users to run parallel programs on cloud using key-value store without editing a single line of code.

Moreover, the load and store instructions translation tool provide a way to run programs on the systems which do not have enough main memory for these programs. By translating the loads and stores to get and put methods, the system will use Bigtable as its main memory source. It is expected that the instrumented program is going to be magnitudes of order slower than the original program. This issue is not covered in this paper.

1.2 Scope

One of the initial goals of the project was to create a tool which can instrument programs written in any user preferred language. Unfortunately, this turned out to be infeasible and not practical in the time span of the project. Communication with Google data store needs a different gRPC library for each language used. Moreover, different LLVM frontends are needed to translate from source language to LLVM Intermediate Representation (on which the actual translation is applied). Currently, LLVM has full support for C and C++ source languages through Clang, while other language frontends have been written using LLVM by the community. As the main aim of the project is to use Bigtable as the key-value store, the use of different source languages seems more like a nice-to-have feature, rather than the essential part of the project. Thus, I decided to make a proof of concept tool for C++ programming language, which can later be extended to instrumenting other source languages.

1.3 Contributions

The contributions of this paper are as follows:

- Research on Google data stores, namely Bigtable, Datastore and Spanner, their features and the consistency models they provide.
- Benchmarked the above data stores based on their throughput and latency using YCSB tool.
- Research on available tools to create calls to Google Bigtable data store (gRPC, protobuf, OpenSHMEM, Intel PIN, LLVM).
- Implemented an LLVM pass which translates all load and store instructions to get and put calls on Bigtable.
- Implemented custom heap allocation functions to minimise main memory usage on instrumented programs.
- Implemented an LLVM pass which translates only load and store instructions originating from dynamic (heap) memory.

1.4 Synopsis

Chapter 2 presents the main requirements for the data stores to be used as distributed shared memory systems. The chapter continues with the background information on the selected Google Cloud data stores, namely Bigtable, Datastore and Spanner. Finally, the chapter discusses the results of the benchmark ran on these data stores.

Chapter 3 starts with the architecture of the tool, also briefly introducing gRPC and protobuf libraries. Then, the chapter briefly talks about the unsuccessful attempt to

translate store and load instructions to get and put operations on data store using Intel PIN tool. The chapter continues with an LLVM pass implementation.

Chapter 4 presents the memory wasting problem, introduced by storing heap variables on Bigtable, and describes the solution - the implementation of custom heap memory allocation functions.

Chapter 5 discusses the correctness and efficiency of the system.

Chapter 6 introduces the API which lets computers on two different locations in the world use the key-value store as distributed shared memory in scenarios like producer/consumer.

Chapter 7 summarizes the work done and possible ways of improving the system.

Chapter 2

Data store for DSM

2.1 Overview

In order to build and test the translation tool, a single cloud data store was chosen to be used as a distributed shared memory system for the project. The main requirements for the data store were:

- provide efficient throughput and latency results;
- provide strong consistency model;
- have a way to run user programs on the same data centre, the data store is located on;
- provide an API to communicate in C++;
- preferably provide key-value database model.

Three Google cloud storages, which met almost all of the requirements, were suggested, namely Bigtable, Datastore and Spanner. Even though neither of the three candidates had key-value store as their primary database model, they were one of the few that provided communication between C++ program and a data store. Google cloud products provide this functionality through gRPC (open source remote procedure call system) using protobuf library and Google APIs. Moreover, all of these Google data storages can be chosen to be located in the same data centre for best throughput and latency results. Further sections provide a brief look into each of the candidates and show the results of the benchmarking on throughput and latency.

2.2 Bigtable

Bigtable (Google, 2018a) is high performance, wide column NoSQL database, which stores data in massively scalable tables, each of which is a sorted key/value map. Tables

consists of rows, each of which is essentially a collection of key/value entries, where the key is a combination of the column family, column qualifier and timestamp.

Bigtable treats all data as raw byte strings. If a row does not include a value for a specific key, the key/value pair simply does not exist. Changes to a row take up extra storage space, as Bigtable stores mutations sequentially and compacts them only periodically, but as the usual amount of data sent from our tool does not exceed 32/64 bits (depending on the machine architecture) the additional amount of memory used is insignificant.

Most importantly, Bigtable supports look up value associated with key operation and provides strong consistency - all writes are seen in the same order.

2.3 Datastore

Datastore (Google, 2018b) is highly-scalable NoSQL, document store model database developed by Google. Unlike Bigtable, it provides a SQL-like query language (GQL) and ACID (Atomicity, Consistency, Isolation, Durability) properties for atomic transactions. Moreover, it supports a variety of data types, including integers, floating-point numbers and many more, although such functionality is not needed for purpose of the project as the tool stores binary data directly. Datastore uses synchronous replication, meaning that data is written to primary storage and the replica simultaneously.

Similarly to Bigtable, Datastore provides strong consistency for entity (row) lookups by key. It also provides strong consistency for ancestor queries but they are not relevant to the project.

2.4 Spanner

Spanner (Google, 2018c) is a horizontally scalable, globally consistent relational database service. Unlike the previously discussed storages, Spanner has an key-value store as additional database model, data scheme and uses SQL. Similarly to the Datastore, it provides ACID transaction properties.

Spanner provides even stronger consistency property than strong consistency, namely external consistency. External consistency guarantees that for any two transactions, T_1 and T_2 : if T_2 starts to commit after T_1 finishes committing, then the timestamp for T_2 is greater than the timestamp for T_1 .

2.5 Benchmarking results

For the benchmarking an existing industry tool was used - Yahoo! Cloud Serving Benchmark (YCSB) (Cooper et al., 2010). A key feature of YCSB, as described by its

developers, is that it is extensible. YCSB is open-source, supports easy definitions of new systems and workloads. Workloads allow to understand the performance tradeoffs of different systems.

The main operations done by the translation tool are reads and writes with a small amount of read-modify-write operations on heap allocation pointer, keeping track of the address to next free memory space. Thus, workloads A and F were chosen, simulating update heavy and read-modify-write using systems, respectively.

For the best results the benchmarking was run on Google Compute Engine (GCE) virtual machine situated at the same data centre as the data stores.

2.5.1 Loading the data

Before running the benchmark on workloads, 1000 rows were inserted into each data store. Figures 2.1 and 2.2 show the latency and throughput achieved by each cloud storage. The results show that both Bigtable and Spanner have much lower latency and higher throughput than Datastore. This can be explained by research results on Datastore using synchronous replication, which makes the host wait until all replications are created, as described in Margaret’s Rouse article Synchronous replication (Rouse, 2016).

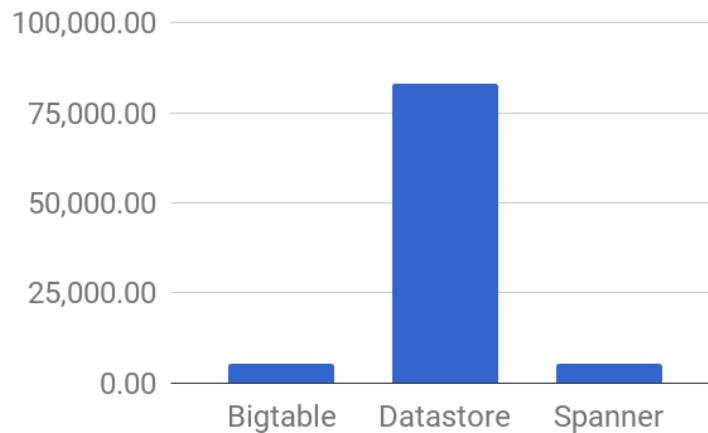


Figure 2.1: Latency (μs) for 1000 insert (write) operations

2.5.2 Workloads

Workload A consists of 1000 operations (500 reads and 500 writes) while workload F consists of 2000 operations (1000 reads, 500 atomic read-modify-write operations and 500 writes). The results of the benchmark in terms of latency on write operations were consistent with the previous loading benchmark results, with Bigtable and Spanner performing significantly better over Datastore (Figure 2.3). The latency on write operations showed a clear dominance by Bigtable.

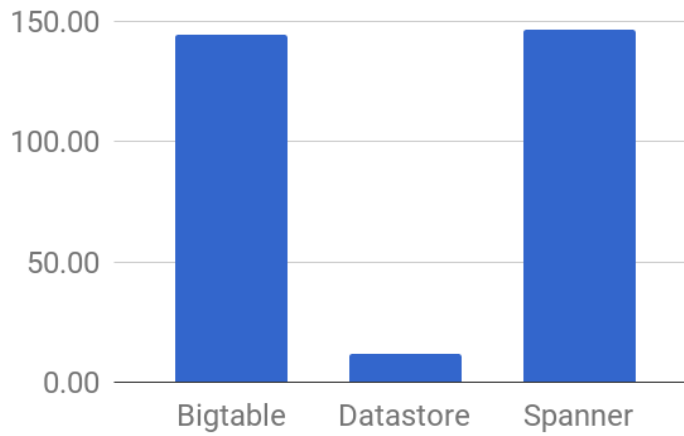


Figure 2.2: Throughput (*ops/sec*) for 1000 insert (write) operations

Even though, the difference on read operations latency between Datastore and two other data storages were smaller than with write operations (Figure 2.4), Datastore still was more than two times slower than Spanner and more than 4 times slower than Bigtable. The latency results on read-modify-write operations showed a similar trend as read and write operations (Figure 2.5).

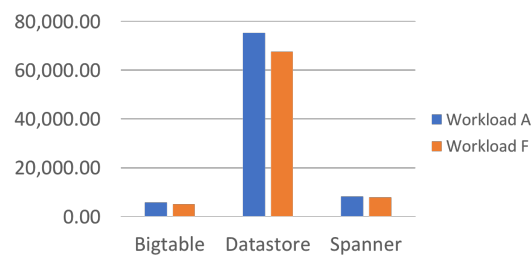


Figure 2.3

The overall throughput, again, showed a significant superiority by Bigtable, as indicated in Figure 2.6.

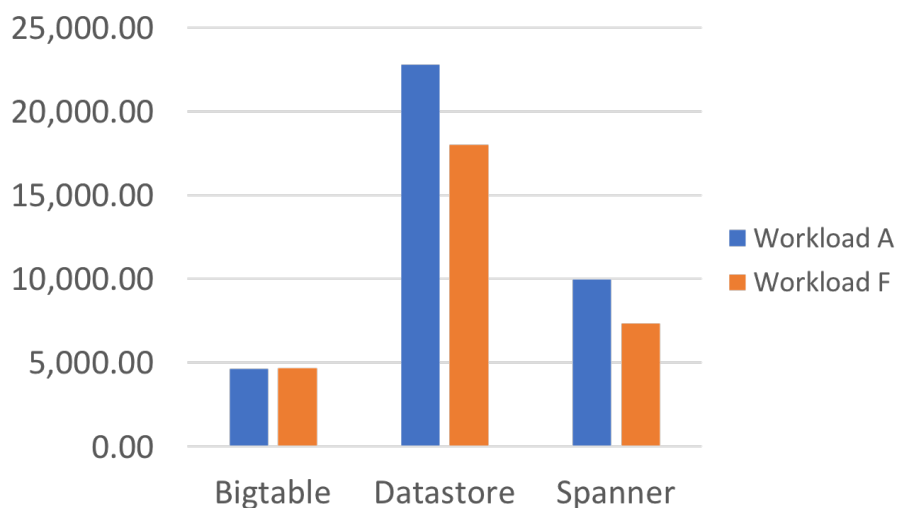
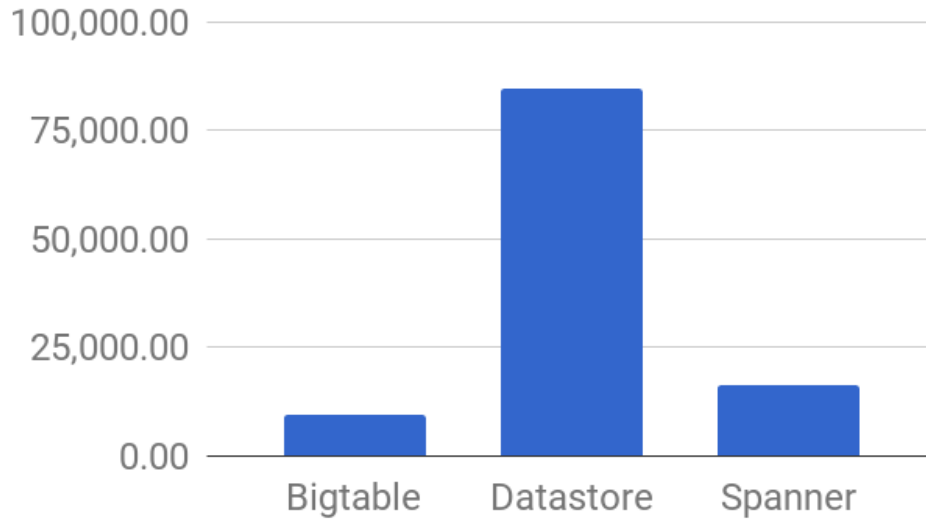
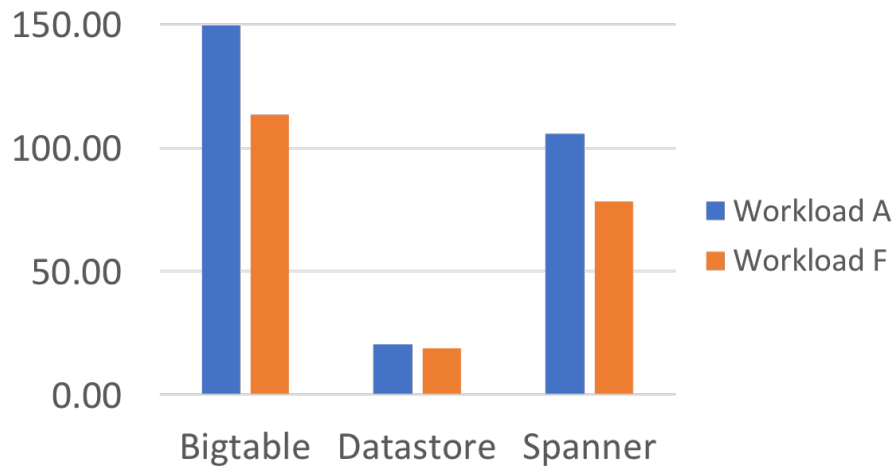


Figure 2.4: Read operations latency (μs) for Workload A (500 reads) and Workload B (1000 reads)

Figure 2.5: Read-modify-operations latency (μs)Figure 2.6: Throughput (*ops/sec*) for Workloads A (1000 operations) and Workload B (2000 operations)

2.5.3 Conclusions

As Bigtable showed the best results in loading of data and on both of the workloads the benchmarks were run on, and since it provided a strong consistency model, it was selected to be used as a distributed shared memory system for the translation tool.

2.6 Reading and writing the contents of Bigtable

Bigtable uses gRPC client to read and write content using C++ remote function calls. According to gRPC webpage (Google, 2018d), gRPC client application can directly

call methods on a server application on a different machine as if it was a local object (similarly to Java RMI). By default gRPC uses protocol buffers, Google's open source language and platform neutral mechanism for serialising structured data. The figure 2.7 visualises communication between gRPC server and client (stub). Bigtable gRPC client is provided through Google APIs repository. In order to simplify the communication with Bigtable table, two functions for putting and getting data to table were implemented, namely `put()` and `get()`.

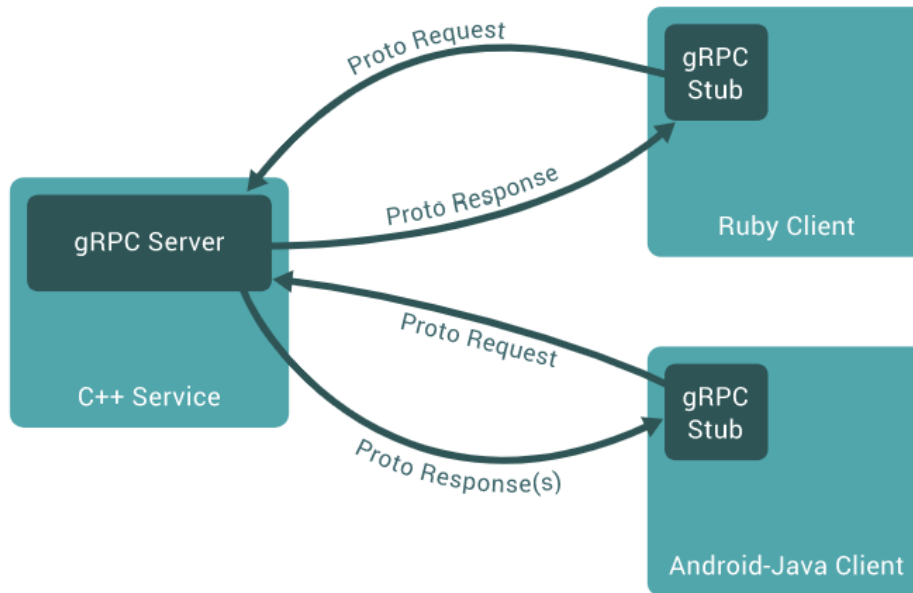


Figure 2.7: Communication between gRPC server and gRPC stubs (clients) (source: Google, 2018d, Guides page)

Put function takes two arguments, 64-bit integer as an address (or key) and 64-bit integer as a value, and does not return anything (see Listing 2.1). First, the arguments are casted to string type. Then a new row mutation request is built, by providing full path to the table (including project, instance and table names), row key, family name and column qualifier and value. Family name and column qualifier are constant as we are using Bigtable as key-value store, thus only one column family and qualifier is used. Finally, the row mutation function is called remotely through Bigtable stub and some status information is stored for debugging purposes.

```

void put(unsigned long long addr, long long val) {
    // cast arguments to string type
    string address = std::to_string(addr);
    string value = std::to_string(val);

    // setup the request
    MutateRowRequest req;
    req.set_table_name(tableName);
    req.set_row_key(address);
    auto setCell = req.add_mutations()->mutable_set_cell();
    setCell->set_family_name(familyName);
    setCell->set_column_qualifier(columnQualifier);
}

```



```

setCell→set_value ( value );

// invoke row mutation on Bigtable
MutateRowResponse resp;
grpc::ClientContext clientContext;
auto status = bigtableStub→MutateRow(&clientContext , req , &resp);
}

```

Listing 2.1: Writing content to Bigtable using put() function

Get function takes an address (or key) with type 64-bit integer as an argument and returns a 64-bit integer value (see Listing 2.2). Similarly to put function, the address value is cast to string. A read row request is created by providing the same full path to the table mentioned above and address string is passed as a row key. The call on rows reading function returns a stream, which is read by chunks and appended to the valueStr variable. As all keys in key-value store are assumed to be unique, the nested loops should run at most one time. Before the value is returned, an if statement checks if the given key had the corresponding value in the table and if so, casts the value to 64-bit integer. If no value was found with corresponding key, the function returns 0.

```

long long get(unsigned long long addr) {
    // convert argument to string type
    string address = to_string(addr);

    // setup the request
    ReadRowsRequest req;
    req.set_table_name(tableName);
    req.mutable_rows()→add_row_keys(address);

    string valueStr;

    // invoke row reading on Bigtable
    auto stream = bigtableStub→ReadRows(&clientContext , req);
    while (stream→Read(&resp)) {
        for (auto& cellChunk : *resp.mutable_chunks()) {
            if (cellChunk.value_size() > 0) {
                valueStr.reserve(cellChunk.value_size());
            }
            valueStr.append(cellChunk.value());
        }
    }

    // convert value to 64-bit integer
    long long value = 0;
    if (!valueStr.empty())
        value = stoll(valueStr);
    }
    return value;
}

```

Listing 2.2: Reading content from Bigtable using get() function

2.7 Issues encountered

Even though both gRPC and protobuf (protocol buffers) libraries are developed by Google, some difficulties were encountered while compiling source builds. The errors were made known to the developers (grp, 2017), but it has slightly stalled the development of the project.

Chapter 3

Calling get/put functions on Bigtable

3.1 Research on possible solutions

Having benchmarked and chosen the cloud data store, the next step was to find out ways to invoke write and read operations on Bigtable with data meant to be shared. Three different strategies were identified: library implementation (or source code level solution), binary translation and compiler level solution.

3.1.1 Source code level solutions

The first option for source code level solution would involve creating an application programming interface (API), which includes calls to Bigtable gRPC client with different types of data. This looked like the easiest solution but it would have introduced the requirement for the user to change the source code in order for the tool to work. Moreover, different data types would require different get and put function overloads, thus the search continued on more generic solutions.

Another source code level solution was considered which involved using OpenSHMEM (Chapman et al., 2010). OpenSHMEM is an open-source partitioned global address space (PGAS) library interface specification. OpenSHMEM implements PGAS by defining remotely accessible data objects (or symmetric data objects) as mechanisms to share information among OpenSMEM processes (also called processing elements) and data objects that are private to each processing element (PE). The interface provides methods to start the OpenSHMEM processing elements in parallel, and communication and synchronization interfaces to access remotely accessible data objects across PEs. The solution using OpenSHMEM involved modifying the API calls for symmetric data objects, changing the storing of shared objects from the symmetric heap (see Figure 3.1) to Bigtable. Even though this solution would introduce some changes to the user source code, OpenSHMEM is a well-known, broadly used API thus more support would be provided for the user than in the first solution. Unfortunately, OpenSHMEM is more targeted for supercomputers or large high-performance

computing clusters with special network setup (i.e. Infiniband network adapters, Qs-Net interconnect, etc), which is not provided by Google's Compute Engine virtual machines (VMs).

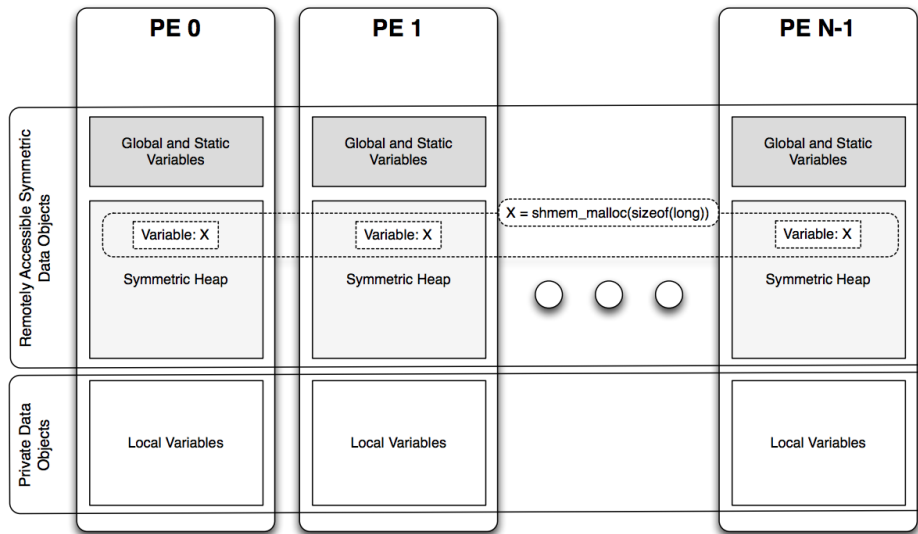


Figure 3.1: OpenSHMEM Memory Model (source: OpenSHMEM-1.3.pdf, 2016)

3.1.2 Compiler level solution

One way to invoke write and read operations on Bigtable on a compiler level is to customise or create a new compiler, which instead of using load and store instructions for shared memory would call get and put function calls for the Bigtable. Of course, creating a new, industry standard compiler would be too excessive and not feasible in the time span of the project, thus only the customisation of already existing compiler was considered. During the research on this option, a solution using LLVM was discovered.

3.2 LLVM

LLVM (Low Level Virtual Machine) is an open source compiler framework for building tools began at the University of Illinois. It supports life-long program analysis and transformation for arbitrary programs. It has an industrial standard compiler (clang/clang++), which has an option to compile C/C++ code to an extensible, strongly typed intermediate representation, namely LLVM IR.

LLVM optimizing compiler, like other industry standard compilers, consists of several components: frontend, optimiser, backend and linker. The important advantage over other compilers is the use of LLVM Bitcode, which consists of a bitstream container format and an encoding of LLVM IR. It provides a clean API boundary separating the compiler frontend and backend, thus making it easier to swap in new frontend and

backend components (Figure 3.2). This is especially useful when developing a new language, as one only needs to create a new frontend component of the compiler and use the provided LLVM optimiser and backend components.

Moreover, as LLVM project is open-source, users can create their own optimisation or transformation passes. Thus it is possible to create an transformation pass for the optimiser, which would iterate over all of the RISC-like instruction set and translate the appropriate load and store instructions to get and put calls for Bigtable.

LLVM has lots of support and documentation on the Internet. Lots of well-known compilers, like NVIDIA's CUDA Compiler or Microsoft DirectX shader compiler, are based on LLVM. All of these features make LLVM a desirable framework to use for our translation tool.

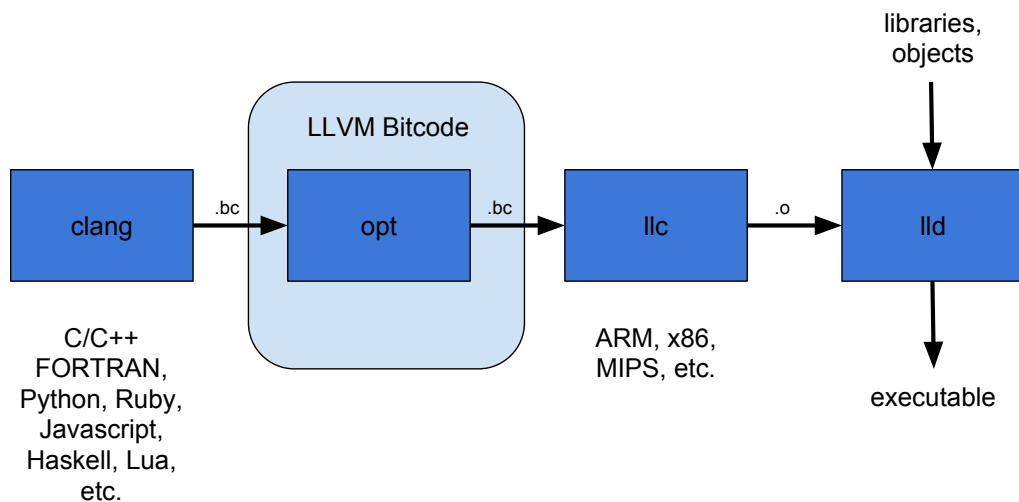


Figure 3.2: LLVM compiler architecture

3.3 Architecture

The translation pass tool consists of 4 transitions (see Figure 3.3). At first, user source code and Bigtable GET/PUT functions are compiled to LLVM IR and combined into a single LLVM Bitcode file. Additional C++ files can be added at this stage if needed later for the LLVM translation pass. After that the code is instrumented by our transformation pass and emits translated LLVM IR into a single LLVM Bitcode file. Then the output is compiled into native machine code and finally linked with gRPC, protobuf and any other relevant libraries to build an executable.

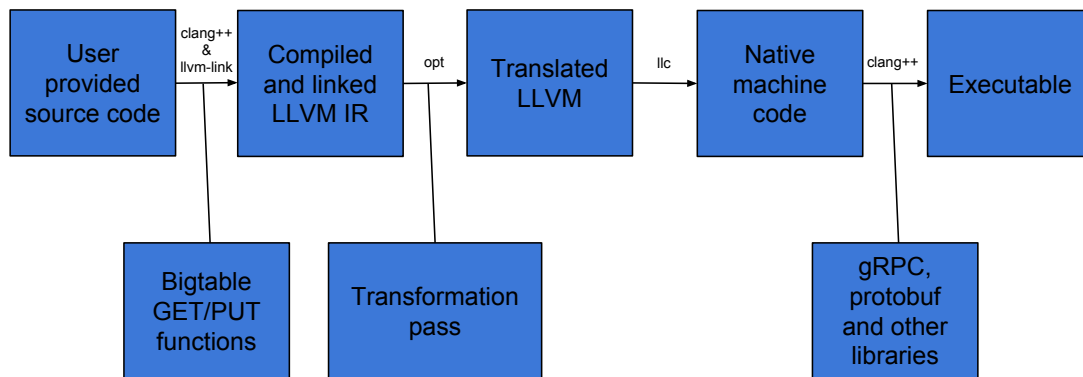


Figure 3.3: LLVM solution architecture

After the preprocessing and compilation of Bigtable GET/PUT functions file, the resulting LLVM Bitcode consists of dozens of gRPC library functions. As we do not want to translate the gRPC load and store instructions, a check was added to the transformation pass, which stops the transformation when the first gRPC function is detected through the iteration. In order for this to work properly, the linking of LLVM Bitcode files in first stage must be done in a strict order: Bigtable GET/PUT functions file must appear after the code that must be translated. This creates a barrier (see Figure 3.4) between the code that is translated and the code which contains functions to be called translated code (i.e. put and get).

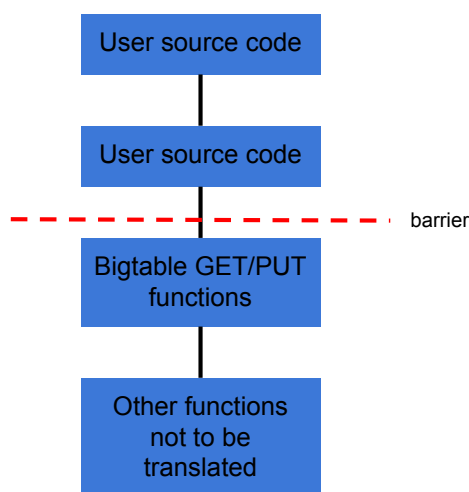


Figure 3.4: Translation barrier

3.4 LLVM pass

The previously mentioned barrier dividing instrumented and uninstrumented code is implemented by checking if the currently iterated function is from Bigtable GET/PUT functions file. The translation pass iterates over all instructions until the barrier. When the iteration reaches this function, the loop is ended. Moreover, the pass also skips over inline functions with external linkage. In LLVM, these functions can be detected by `hasLinkOnceODRLinkage` function call.

Load and store instructions accept different types of arguments (remember IR is a strongly typed instruction set). For store instructions, the address must have a higher pointer indirection degree than the value. For instance, the corresponding store instruction for `int * x = &y` (assuming `y` is of type integer) would have the address with the type of pointer to pointer to integer and the value with the type of pointer to integer. For load function translations, get function must return the value with the same type as load instruction. Moreover, get and load functions only accepts 64-bit integer arguments, thus sign extension instructions might be needed. This makes the translation of store (and get) instructions more complicated than just replacing one instructions with the others.

Store to put instruction translation starts by converting the address pointer to 64-bit integer with `ptrtoint` instruction. If the value is of integer type and not 64 bits wide it is sign extended using `sext` instruction. If the value is of pointer type it is cast to 64-bit integer by `ptrtoint` instruction. Otherwise, it is assumed to be 64-bit integer¹. Finally, both arguments now being of 64-bit² integer type are given as arguments to `put` function call. The figure 3.5 shows a subset of instruction set before and after store instruction translation.

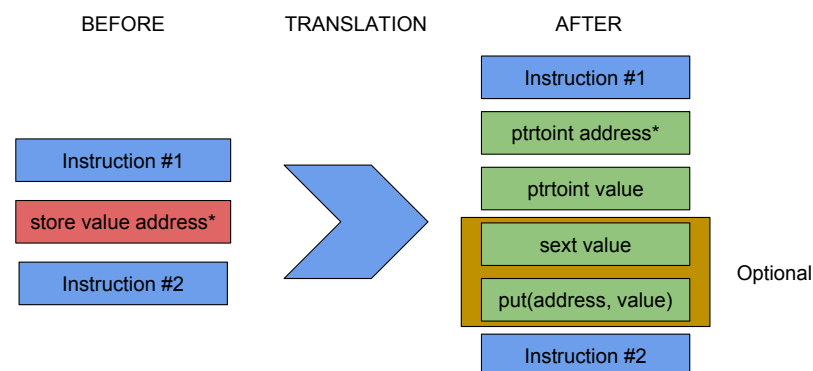


Figure 3.5: LLVM instruction set before and after store instruction translation

¹ Although it might of other type (i.e. struct).

² 64-bit integer type for get and put instruction arguments and/or return type was chosen deliberately. This is the maximum biggest type of integer, thus it makes the casting part a bit simpler. For example, if 32-bit integer type was chosen, some values would need to be sign extended while others would need to be truncated.

Load instruction translation begins by identifying its return type and pointer indirection degree (only relevant if of pointer type). Similarly to store translation, the address is cast to 64-bit integer using `ptrtoint` instruction, and passed to `get` as an argument. As the return type of `get` function is 64-bit integer, it must be cast to the appropriate type (unless the load instruction actually returns 64-bit integer). If the returned type is integer it is truncated to the expected integer type using `trunc` instruction. Finally, if the expected returned type is pointer, the resulting value is converted to a pointer type with an appropriate pointer indirection degree (identified at the beginning) with `inttoptr` instruction. The figure 3.6 shows a subset of instruction set before and after store instruction translation.

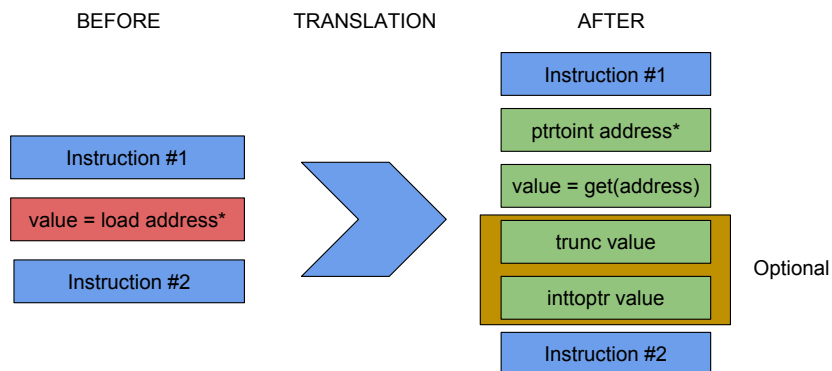


Figure 3.6: LLVM instruction set before and after get instruction translation

Chapter 4

Custom heap allocator

4.1 Motivation

Even though all load and store instructions were translated using the translation pass, the space for heap was still being allocated on the main memory. This was detected using Valgrind, a tool for memory debugging and profiling. The issue was solved by developing a custom first-fit free-list heap allocator, which does not allocate space from main memory.

4.2 Implementation

The heap allocator implementation was based on Marwan Burelle's malloc tutorial (Burelle, 2009) and adjusted to work on Bigtable. The heap allocator implements four functions: malloc, free, realloc and calloc. All of these functions are counterparts of the Standard C++ Library functions and have the same function definitions.

The allocator keeps a list of meta-data blocks, which contain information about chunks of data allocated with malloc. This lets the memory be reused after it has been released with call to free function. Unlike in tutorial implementation, only meta-data objects are stored on main memory. The actual requested space is allocated on Bigtable. Thus, the implemented custom heap allocator uses two heaps: main memory and Bigtable. The figure 4.1 sketches the memory organisation of two heaps. The structure of meta-data blocks described in the tutorial were adjusted to reflect these changes. Besides storing the size of data block and pointers to other blocks, the meta-data blocks also stores the address of data block on Bigtable.

The main memory heap is managed by the default memory allocator (Standard C++ library provided malloc, free, etc). For the custom heap allocator, it is used as meta-data objects storage.

The heap on Bigtable is implemented as a continuous space of memory with two bounds: the start of the heap and the end point called the Bigtable break. The start

of the heap is initialized on the first call on custom malloc function by calling sbrk function with an increment equal to 0 (this returns the break address on the main memory heap). Thus, the addresses of main memory and Bigtable heap are identical. The end of the heap is managed by set_bt_brk function, which was implemented to reflect the behaviour of sbrk function (see Listing 4.1).

```
void* set_bt_brk(int incr) {
    // if uninitialised, set to sbrk(0)
    if (current_bt_break == 0) {
        current_bt_break = (uintptr_t) sbrk(0);
        initial_bt_break = current_bt_break;
    }
    uintptr_t old_break = current_bt_break;
    current_bt_break += incr;
    return (void*) old_break;
}
```

Listing 4.1: set_bt_brk function implementation

Of course, another solution could be to allocate both meta-data objects and the actual requested space on the Bigtable, but this would increase the communication costs with Bigtable. As meta-data objects only take up to 40 bytes (on a 64-bit system), the decision was made to store them locally.

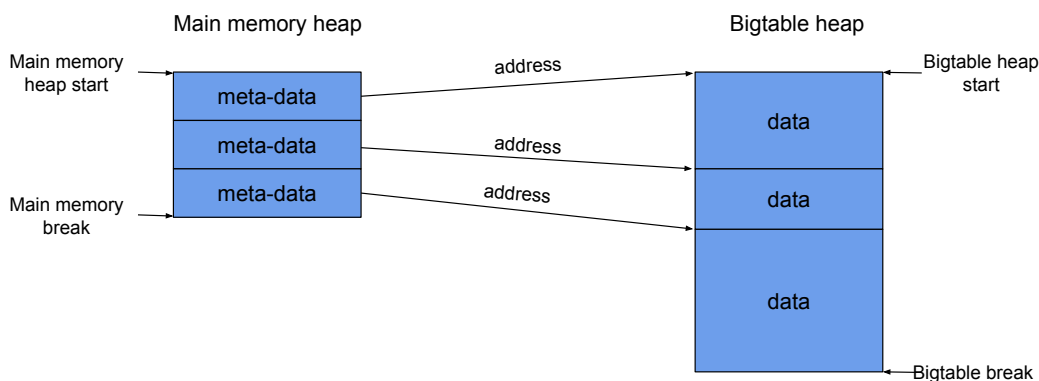


Figure 4.1: LLVM instruction set before and after get instruction translation

Malloc function starts by changing the requested size to be a multiple of 4 to align the pointers by 32 bits. If the heap is not empty (already called malloc() before), the linked list of meta-data blocks is searched for the first free chunk that is wide enough for the request. If such block is found and the difference between the requested size and the size of the block is enough to store a minimum allowed block (32 bytes¹) the block is split into two blocks. The first block is marked as used. If the heap is empty or no existing wide enough block is found, the heap is extended. Finally, the address to the block on Bigtable is returned.

¹In the popular dlmalloc (Lea) implementation, the smallest allowed allocation is 32 bytes on a 64-bit system, thus I decided to use the same size in my malloc implementation.

Free function accepts a pointer to heap memory block to be freed and starts by checking if it points to is a valid Bigtable heap address. If it does, the pointer address is used to retrieve the meta-data block. In order to do this efficiently a hash table was introduced to map the data block address to meta-data block. Every time a heap is extended or a block is split, the hash table is updated with a new key-value (address and meta-data block) pair. This is a more efficient solution than iterating over the linked-list of meta-data blocks. After the meta-data block is retrieved, it is freed. If any of the neighbouring meta-data blocks are free, the blocks are fused into one (to cope with fragmentation). If after the fusion the resulting meta-block is the tail of the linked-list the memory is released. This part was adjusted to decrease the Bigtable break by the size of the block being released. Consequently, the corresponding pair is deleted from the hash table and the meta-data block is freed using default malloc.

Realloc functions accepts two arguments, pointer to existing heap memory block and the new size for the block. If the pointer address is a valid heap address, the meta-data block is fetched using the hash table mentioned above. The size is changed to be aligned with 32-bit pointers. If the requested size is smaller than the original size, the block is split in two. Otherwise, if the next block is free and provide enough space (combined with the original block), the two blocks are fused into one and split if necessary. If none of the options above are true, a new block is allocated with malloc, the old block contents are copied to the new one and, finally, the old block is freed. The block copying procedure was modified to work with Bigtable. The old block values are copied using a combination of GET and PUT function calls to Bigtable (see Listing 4.2). Lastly, if the pointer given as an argument to realloc is null, the behaviour is the same as calling malloc with the given size.

```
void copy_block(block src , block dst) {
    int *sdata , *ddata;
    unsigned long long value , *a , b;
    sdata = (int*) src->addr;
    ddata = (int*) dst->addr;
    for (size_t i = 0; i*4 < src->size && i*4 < dst->size; i++) {
        // convert int* to 64-bit integer
        a = (unsigned long long*) &sdata[i];
        b = (unsigned long long) *a;
        value = get(b);

        // convert int* to 64-bit integer
        a = (unsigned long long*) &ddata[i];
        b = (unsigned long long) *a;
        put(b, value);
    }
}
```

Listing 4.2: copy_block implementation

Calloc function accepts an integer representing a number of elements to allocate and an integer representing the size of each element. First, the malloc is called with the product of two arguments. The new block is iterated by 32 bit steps and initialised with 0 values. Again, the implementation was modified to work with Bigtable, using PUT function (see Listing 4.3).

```
s4 = align4(num * size) >> 2;
for (i = 0; i < s4; i++) {
    // convert int* to 64-bit integer
    unsigned long long* a = (unsigned long long*) &new_block[i];
    unsigned long long b = (unsigned long long) a;
    put(b, 0ULL);
}
```

Listing 4.3: new_block initialisation with zeroes

4.3 Consistency on multithreaded programs

Even though the malloc tutorial was very helpful in implementing a custom heap allocator, it didn't mention anything about allocator consistency on multithreaded programs. The simplest solution to this problem was implemented using a lock. A single mutex was created and the calls to its functions lock and unlock were added to the entry and exit points of malloc, free, realloc and calloc functions. This synchronises the above functions and lets only one thread to make modifications to the heap. This means that no other threads are permitted to do allocations and releases while the other thread is modifying the heap by either of the above functions. Even though this is a correct solution, it is very inefficient. Threads that frequently allocate and release memory from the heap are constantly being blocked or block others threads, essentially making the execution of the program serial.

Another way to solve the thread-safety problem is by allocating a large chunk of memory off the heap to each thread and then managing the space within the thread. However, some threads might not be using all of their allocated memory, which results in poor memory usage efficiency. Moreover, some threads might need more memory than given by default, thus there should be ways to increase the per thread heap memory. Thus, it can be seen that the increase in performance adds additional complexity to the allocator. Clearly, this is a more involved solution and due to a strict time limits for the project it was not implemented. This is one of the areas where the translation tool could be improve in the future work.

For comparison purposes, it was decided to also implement a heap allocator with only a malloc function, without any memory releasing strategy. Heap memory address is incremented with an atomic read-modify-write operation provided by Bigtable. This way the calls to malloc don't have to be synchronised explicitly, as the job is done by atomic calls to Bigtable.

Chapter 5

Translation improvements

- 5.1 Different type of arguments to get and put functions**
- 5.2 Detecting store and load instructions to heap memory**

Chapter 6

Results

6.1 Benchmarking heap allocators

6.2 Full and heap translation comparison

6.3 Comparison to uninstrumented program

Chapter 7

Conclusions

7.1 Overview

7.2 Future work

7.2.1 Better support for structured data

Bibliography

- grpc c++ fails to compile (zlib), December 2017. URL <https://github.com/grpc/grpc/issues/13640>.
- Emery D. Berger, Benjamin G. Zorn, and Kathryn S. McKinley. Reconsidering custom memory allocation. *SIGPLAN Not.*, 37(11):1–12, November 2002. ISSN 0362-1340. doi: 10.1145/583854.582421. URL <http://doi.acm.org/10.1145/583854.582421>.
- Marwan Burelle. *A Malloc Tutorial*. Laboratoire Systeme et Securite de l’EPITA (LSE), February 2009.
- Barbara Chapman, Tony Curtis, Swaroop Pophale, Stephen Poole, Jeff Kuehn, Chuck Koelbel, and Lauren Smith. Introducing openshmem: Shmem for the pgas community. 10 2010.
- Brian F. Cooper, Adam Silberstein, Erwin Tam, Raghu Ramakrishnan, and Russell Sears. Benchmarking cloud serving systems with ycsb. In *Proceedings of the 1st ACM Symposium on Cloud Computing*, SoCC ’10, pages 143–154, New York, NY, USA, 2010. ACM. URL <https://www2.cs.duke.edu/courses/fall13/cps296.4/838-CloudPapers/ycsb.pdf>.
- Google, 2018a. URL <https://cloud.google.com/bigtable/>.
- Google, 2018b. URL <https://cloud.google.com/datastore/>.
- Google, 2018c. URL <https://cloud.google.com/spanner/>.
- Google, 2018d. URL <https://grpc.io/docs/guides>.
- Katarina Grolinger, Wilson A Higashino, Abhinav Tiwari, and Miriam AM Capretz. Data management in cloud environments: Nosql and newsql data stores. *Journal of Cloud Computing: Advances, Systems and Applications*, December 2013.
- D Lea. Dlmalloc, 2010. URL <http://g.oswego.edu/dl/html/malloc.html>.
- OpenSHMEM-1.3.pdf. Openshmem: Application programming interface, February 2016.
- Margaret Rouse. Synchronous replication. *WhatIs.com*, 2016.
- Z.Huang, W.Chen, M.Purvis, and W.Zheng. Vodca: View-oriented, distributed, cluster-based approach to parallel computing. In *Cluster Computing and the Grid*,

2006. *CCGRID 06. Sixth IEEE International Symposium on*, volume 2, pages 15–15, May 2006.